

Linux Kernel主要通过三类机制来实现SMP（Symmetric Multiprocessing，对称多核）系统CPU core的电源管理：

- cpu hotplug: 根据应用场景来up/down CPU
- cpuidle framework: 当cpu上没有可执行任务时，就会进入空闲状态
- cpufreq framework: 根据使用场景和系统负荷来调整CPU的电压和频率

cpufreq framework的核心功能，是通过调整CPU core的电压或频率，兼顾系统的性能和功耗。在不需要高性能时，降低电压或频率，以降低功耗；在需要高性能时，提高电压或频率，以提高性能。

cpufreq framework中的几个重要概念：

1. policy（策略）：同一个簇的CPU动态调频的一个集合结构体，包含了当前使用的governor和cpufreq driver
2. governor（调节器）：决定如何计算合适的频率或电压
3. cpufreq driver：来实现真正的调频执行工作（与平台相关）

常用的governor类型

1. Performance：总是将CPU置于最高性能的状态，即硬件所支持的最高频率、电压
2. Powersaving：总是将CPU置于最节能的状态，即硬件所支持的最低频率、电压
3. Ondemand：设置CPU负载的阈值T，当负载低于T时，调节至一个刚好能够满足当前负载需求的最低频/最低压；当负载高于T时，立即提升到最高性能状态
4. Conservative：跟Ondemand策略类似，设置CPU负载的阈值T，当负载低于T时，调节至一个刚好能够满足当前负载需求的最低频/最低压；但当负载高于T时，不是立即设置为最高性能状态，而是逐级升高主频/电压
5. Userspace：将控制接口通过sysfs开放给用户，由用户进行自定义策略
6. Schedutil：这是从Linux-4.7版本开始才引入的策略，其原理是根据调度器所提供的CPU利用率信息进行电压/频率调节，EAS使用schedutil进行调频

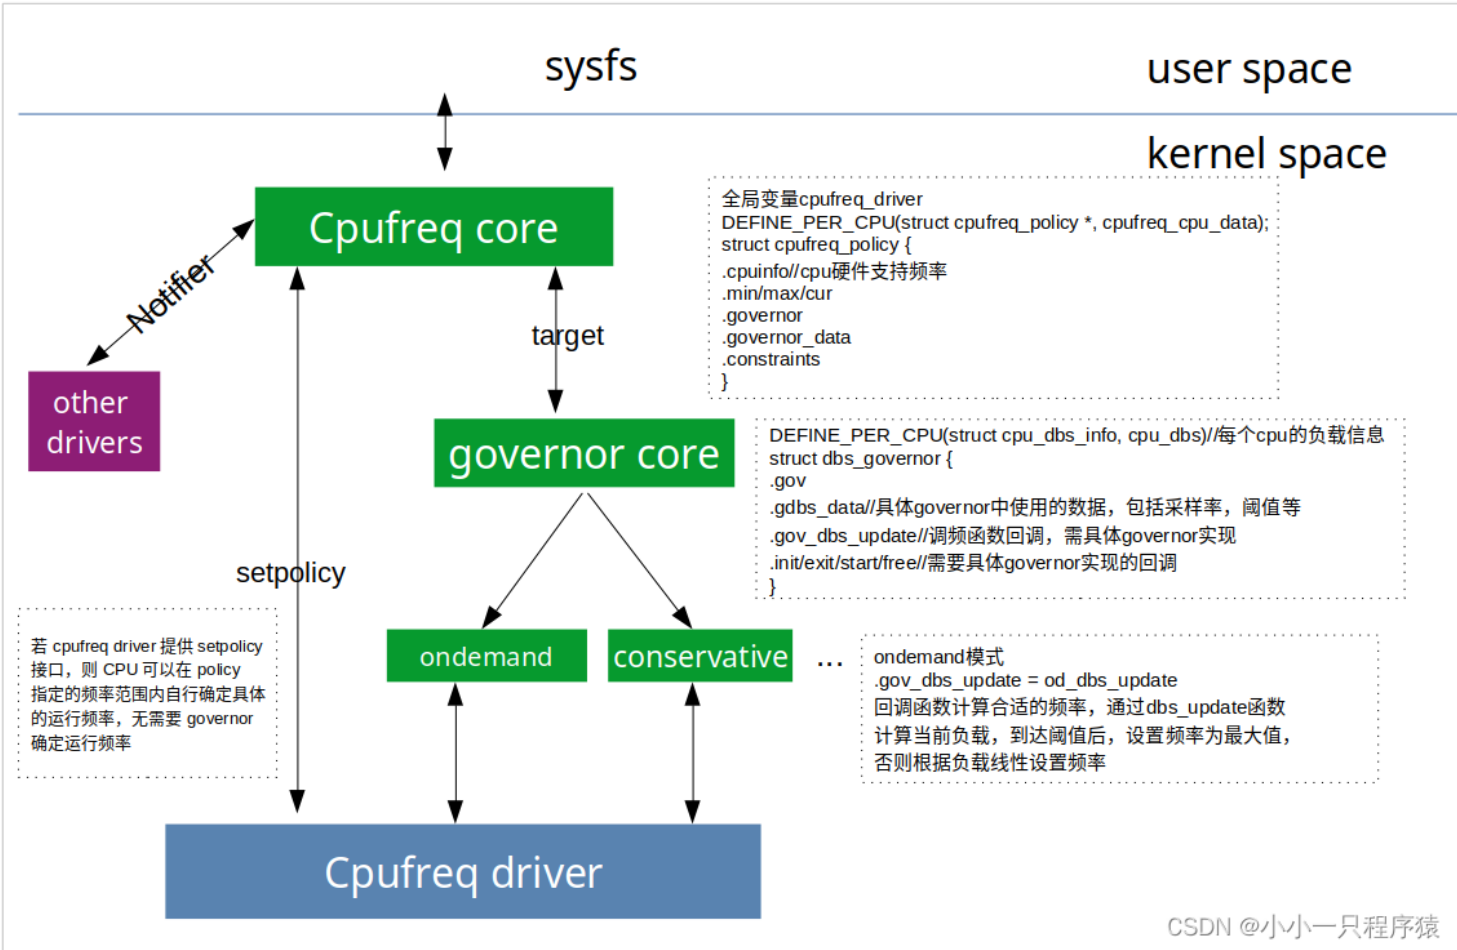
sysfs用户层接口，目录位于 /sys/devices/system/cpu/cpufreq/policy

```
root@Ubuntu:/sys/devices/system/cpu/cpufreq/policy0# ls
affected_cpus          related_cpus           scaling_max_freq
cpuinfo_cur_freq       scaling_available_frequencies scaling_min_freq
cpuinfo_max_freq       scaling_available_governors scaling_setspeed
cpuinfo_min_freq       scaling_cur_freq       stats
cpuinfo_transition_latency scaling_driver
ondemand               scaling_governor
root@Ubuntu:/sys/devices/system/cpu/cpufreq/policy0#
```

名称	说明
cpuinfo_max_freq	硬件所支持的最高频率
cpuinfo_min_freq	硬件所支持的最低频率
affected_cpus	该policy影响到哪些cpu（没显示offline状态的cpu）
related_cpus	该policy影响到的所有cpu，包括offline状态的cpu
scaling_max_freq	该policy支持调整的最高频率
scaling_min_freq	该policy支持调整的最低频率

名称	说明
scaling_cur_freq	policy当前设置的频率
scaling_available_governors	当前系统支持的governor
scaling_available_frequencies	支持的调频频率
scaling_driver	当前使用的调频驱动
scaling_governor	当前使用的governor
scaling_setspeed	在userspace模式下才能使用，手动设置频率

cpufreq软件架构

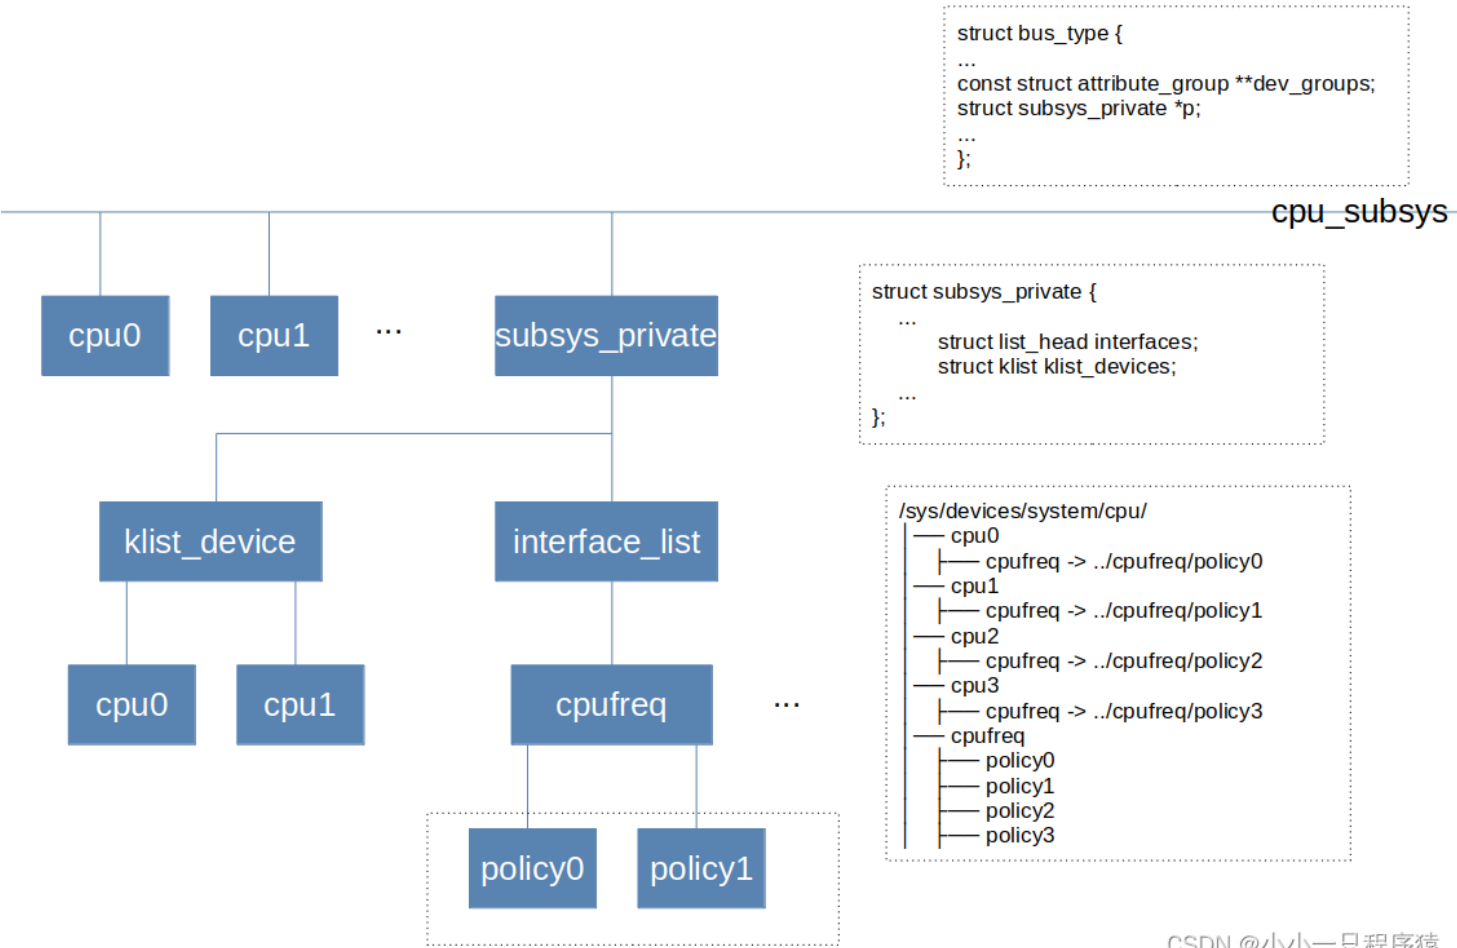


CSDN @小小一只程序猿

cpufreq core（可以理解为对policy的操作）：把一些公共的逻辑和接口代码抽象出来

- cpufreq 作为所有cpu设备的一个功能，注册到了 cpu\_subsys 总线上
- 对上以sysfs的形式向用户空间提供统一的接口，以notifier的形式向其他driver提供频率变化的通知
- 对下提供CPU频率和电压控制的驱动框架，方便底层driver的开发，同时提供governor框架，用于实现不同的频率调整机制
- 内部封装各种逻辑，主要围绕 struct cpufreq\_policy struct cpufreq\_driver struct cpufreq\_governor 三个数据结构进行

kernel使用 struct cpufreq\_policy 用来抽象cpufreq，它从一定程度上代表了一个CPU簇的cpufreq的属性



CSDN @小小一只程序猿

cpufreq\_policy 结构体

```

struct cpufreq_cpuinfo {
    unsigned int      max_freq;           // cpu最大频率
    unsigned int      min_freq;           // cpu最小频率

    /* in 10-9 s = nanoseconds */
    unsigned int      transition_latency; // cpu频率转换时间 单位: ns
};

struct cpufreq_policy {
    /* CPUs sharing clock, require sw coordination */
    cpumask_var_t      cpus; /* Online CPUs only */
    cpumask_var_t      related_cpus; /* Online + Offline CPUs */
    cpumask_var_t      real_cpus; /* Related and present */

    unsigned int        shared_type; /* ACPI: ANY or ALL affected CPUs
                                     should set cpufreq */
    unsigned int        cpu; /* cpu managing this policy, must be online */

    struct clk          *clk;
    struct cpufreq_cpuinfo cpuinfo; /* see above */

    unsigned int        min; /* in kHz */
    unsigned int        max; /* in kHz */
    unsigned int        cur; /* in kHz, only needed if cpufreq
                             * governors are used */

    struct cpufreq_governor *governor; /* see below */
    void                *governor_data;
    char                last_governor[CPUFREQ_NAME_LEN]; /* last governor used */

    struct work_struct  update; /* if update_policy() needs to be
                             * called, but you're in IRQ context */

    struct cpufreq_user_policy user_policy;
    struct cpufreq_frequency_table *freq_table;
    enum cpufreq_table_sorting freq_table_sorted;

    struct list_head    policy_list;
    struct kobject       kobj;
    struct completion    kobj_unregister;

    /*
     * Preferred average time interval between consecutive invocations of
     * the driver to set the frequency for this policy. To be set by the
     * scaling driver (0, which is the default, means no preference).
     */
    unsigned int        transition_delay_us;

    /*
     * Remote DVFS flag (Not added to the driver structure as we don't want
     * to access another structure from scheduler hotpath).
     *
     * Should be set if CPUs can do DVFS on behalf of other CPUs from
     * different cpufreq policies.
     */
    bool                dvfs_possible_from_any_cpu;

    /* Cached frequency lookup from cpufreq_driver_resolve_freq. */
    unsigned int        cached_target_freq;

```

```

    int cached_resolved_idx;

    /* cpufreq-stats */
    struct cpufreq_stats    *stats;

    /* For cpufreq driver's internal use */
    void                    *driver_data;

    ...
};

```

driver/cpufreq/cpufreq.c 中定义了一个全局的percpu变量

```
static DEFINE_PER_CPU(struct cpufreq_policy *, cpufreq_cpu_data);
```

这里对应E2000 sysfs中3个policy文件夹，两个小核在一个簇中，使用1个policy，另外两个大核分别对应1个policy

```

root@Ubuntu:/sys/devices/system/cpu/cpufreq# tree -L 1
.
├── policy0
├── policy2
└── policy3

```

per-CPU变量是linux系统一个非常重要的特性，它为系统中的每个处理器都分配了该变量的副本。这样做的好处是，在多处理器系统中，当处理器操作属于它的变量副本时，不需要考虑与其他处理器的竞争的问题，同时该副本还可以充分利用处理器本地的硬件缓冲cache来提供访问速度

## cpufreq初始化过程

cpufreq\_driver 结构体如下：

```

struct cpufreq_driver {
    char            name[CPUFREQ_NAME_LEN];
    u16             flags;
    void            *driver_data;

    /* needed by all drivers */
    int             (*init)(struct cpufreq_policy *policy);
    int             (*verify)(struct cpufreq_policy_data *policy);
    int             (*target_index)(struct cpufreq_policy *policy,
                                    unsigned int index);
    unsigned int    (*fast_switch)(struct cpufreq_policy *policy,
                                    unsigned int target_freq);

    /* should be defined, if possible */
    unsigned int    (*get)(unsigned int cpu);

    /* Called to update policy limits on firmware notifications. */
    void            (*update_limits)(unsigned int cpu);

    int             (*online)(struct cpufreq_policy *policy);
    int             (*offline)(struct cpufreq_policy *policy);
    int             (*exit)(struct cpufreq_policy *policy);

    struct freq_attr **attr;
};

```

## cpufreq初始化概述

在kconfig中(CPU Power Management -> CPU Frequency scaling)可以对cpufreq进行配置，可以配置支持的governor及系统默认的governor，以及cpufreq调频driver，例如Phytium E2000 5.10内核的配置如下，默认使用schedutil governor，根据调度器所提供的CPU利用率信息进行电压/频率调节，EAS能源感知依赖该governor工作：

```

[*] CPU Frequency scaling
[*] CPU frequency transition statistics
    Default CPUFreq governor (schedutil) --->
    *- 'performance' governor
    <M> 'powersave' governor
    <*> 'userspace' governor for userspace frequency scaling
    <*> 'ondemand' cpufreq policy governor
    <M> 'conservative' cpufreq governor
    *- 'schedutil' cpufreq policy governor
    *** CPU frequency scaling drivers ***
    <*> Generic DT based cpufreq driver
    <M> CPUFreq driver based on the ACPI CPPC spec
    <*> SCPI based CPUfreq driver
    <*> SCMI based CPUfreq driver

```

cpufreq的初始化从cpufreq\_driver注册开始， cpufreq\_register\_driver() 函数为cpufreq驱动注册的入口，驱动程序通过调用该函数进行初始化，传入相关的 struct cpufreq\_driver ， cpufreq\_register\_driver() 会调用 subsys\_interface\_register() 最终执行回调函数 cpufreq\_add\_dev ， 然后调用 cpufreq\_online() 走初始化流程

## Performance Domain opp (Operating Performance Points) 表初始化

OPP表的定义：域中每个设备支持的电压和频率的离散元组的集合称为Operating Performance Points (OPP) ,内核设备树opp文档 Documentation/devicetree/bindings/opp/opp.txt

假设一个CPU设备支持如下的电压和频率关系：

{300MHz at minimum voltage of 1V}

{800MHz at minimum voltage of 1.2V}

{1GHz at minimum voltage of 1.3V}

用OPP表示就可以用{Hz, uV}方式表示如下：

{300000000, 1000000}

{800000000, 1200000}

{1000000000, 1300000}

这里初始化的就是各个性能域（即不同CPU簇）的OPP表，在E2000平台中是通过SCMI的Performace domain management protocol协议获取PERFORMANCE\_DESCRIBE\_LEVELS这个参数表，具体的协议实现源码

在 `drivers/firmware/arm_scmi/perf.c` 里面， `perf.c` 实现了SCMI的Performance domain managment protocol，`scmi cpufreq_drvier`也是通过 `perf_ops` 函数集进行调频

```

// include/linux/scmi_protocol.h
// 抽象描述scmi协议的结构体，相应的ops操作集对应scmi的一个协议
struct scmi_handle {
    struct device *dev;
    struct scmi_revision_info *version;
    const struct scmi_perf_ops *perf_ops;
    const struct scmi_clk_ops *clk_ops;
    const struct scmi_power_ops *power_ops;
    const struct scmi_sensor_ops *sensor_ops;
    const struct scmi_reset_ops *reset_ops;
    const struct scmi_notify_ops *notify_ops;
    /* for protocol internal use */
    void *perf_priv;
    void *clk_priv;
    void *power_priv;
    void *sensor_priv;
    void *reset_priv;
    void *notify_priv;
    void *system_priv;
};

// scmi_handle这个结构体实现的就是scmi整个协议的处理
static const struct scmi_handle *handle;

// include/linux/scmi_protocol.h
/**
 * struct scmi_perf_ops - represents the various operations provided
 * by SCMI Performance Protocol
 *
 * @limits_set: sets limits on the performance level of a domain
 * @limits_get: gets limits on the performance level of a domain
 * @level_set: sets the performance level of a domain
 * @level_get: gets the performance level of a domain
 * @device_domain_id: gets the scmi domain id for a given device
 * @transition_latency_get: gets the DVFS transition latency for a given device
 * @device_ops_add: adds all the OPPs for a given device
 * @freq_set: sets the frequency for a given device using sustained frequency
 * to sustained performance level mapping
 * @freq_get: gets the frequency for a given device using sustained frequency
 * to sustained performance level mapping
 * @est_power_get: gets the estimated power cost for a given performance domain
 * at a given frequency
 */
struct scmi_perf_ops {
    int (*limits_set)(const struct scmi_handle *handle, u32 domain,
                      u32 max_perf, u32 min_perf);
    int (*limits_get)(const struct scmi_handle *handle, u32 domain,
                      u32 *max_perf, u32 *min_perf);
    int (*level_set)(const struct scmi_handle *handle, u32 domain,
                     u32 level, bool poll);
    int (*level_get)(const struct scmi_handle *handle, u32 domain,
                     u32 *level, bool poll);
    int (*device_domain_id)(struct device *dev);
    int (*transition_latency_get)(const struct scmi_handle *handle,
                                  struct device *dev);
    int (*device_ops_add)(const struct scmi_handle *handle,
                           struct device *dev);
    int (*freq_set)(const struct scmi_handle *handle, u32 domain,
                    unsigned long rate, bool poll);

```



```

    int (*freq_get)(const struct scmi_handle *handle, u32 domain,
                    unsigned long *rate, bool poll);
    int (*est_power_get)(const struct scmi_handle *handle, u32 domain,
                        unsigned long *rate, unsigned long *power);
    bool (*fast_switch_possible)(const struct scmi_handle *handle,
                                struct device *dev);
};

// scmi performance domain management protocol(性能域管理相关协议 0x13), messageid: 0x03
// scmi opp结构体定义
struct scmi_opp {
    u32 perf;           // 性能级别, 单位KHz
    u32 power;          // 当前性能级别的功耗
    u32 trans_latency_us; // 切换延时
};

// scmi performance domain management protocol(性能域管理相关协议 0x13)对应操作函数集
// scmi cpufreq_driver 主要利用这个函数集进行调频相关操作
// 对应Performance domain management protocol各个message_id
static const struct scmi_perf_ops perf_ops = {
    .limits_set = scmi_perf_limits_set,
    .limits_get = scmi_perf_limits_get,
    .level_set = scmi_perf_level_set,
    .level_get = scmi_perf_level_get,
    .device_domain_id = scmi_dev_domain_id,
    .transition_latency_get = scmi_dvfs_transition_latency_get,
    .device_opps_add = scmi_dvfs_device_opps_add,
    .freq_set = scmi_dvfs_freq_set,
    .freq_get = scmi_dvfs_freq_get,
    .est_power_get = scmi_dvfs_est_power_get,
    .fast_switch_possible = scmi_fast_switch_possible,
};

// 在这个宏进行SCMI performance domain management protocol协议的初始化
DEFINE_SCMI_PROTOCOL_REGISTER_UNREGISTER(SCMI_PROTOCOL_PERF, perf)

#define DEFINE_SCMI_PROTOCOL_REGISTER_UNREGISTER(id, name) \
int __init scmi_##name##_register(void) \
{ \
    return scmi_protocol_register((id), &scmi_##name##_protocol_init); \
} \
\
void __exit scmi_##name##_unregister(void) \
{ \
    scmi_protocol_unregister((id)); \
}
// 展开该宏
int __init scmi_perf_register(void)
{
    return scmi_protocol_register(SCMI_PROTOCOL_PERF, &scmi_perf_protocol_init);
}

// 初始化过程中调用了scmi_perf_protocol_init();
static int scmi_perf_protocol_init(struct scmi_handle *handle)
{
    int domain;
    u32 version;
    struct scmi_perf_info *pinfo;

```

```

// 获取当前perf domain management协议版本
scmi_version_get(handle, SCMI_PROTOCOL_PERF, &version);

dev_dbg(handle->dev, "Performance Version %d.%d\n",
        PROTOCOL_REV_MAJOR(version), PROTOCOL_REV_MINOR(version));

pinfo = devm_kzalloc(handle->dev, sizeof(*pinfo), GFP_KERNEL);
if (!pinfo)
    return -ENOMEM;

scmi_perf_attributes_get(handle, pinfo);

pinfo->dom_info = devm_kcalloc(handle->dev, pinfo->num_domains,
                              sizeof(*pinfo->dom_info), GFP_KERNEL);
if (!pinfo->dom_info)
    return -ENOMEM;

// 遍历每个performance_domain, 获取performance domain的属性和performance level参数
for (domain = 0; domain < pinfo->num_domains; domain++) {
    struct perf_dom_info *dom = pinfo->dom_info + domain;

    // 获取performance domain属性
    scmi_perf_domain_attributes_get(handle, domain, dom);
    // 获取performance level参数即opp表
    scmi_perf_describe_levels_get(handle, domain, dom);

    if (dom->perf_fastchannels)
        scmi_perf_domain_init_fc(handle, domain, &dom->fc_info);
}

scmi_register_protocol_events(handle,
                              SCMI_PROTOCOL_PERF, SCMI_PROTO_QUEUE_SZ,
                              &perf_event_ops, perf_events,
                              ARRAY_SIZE(perf_events),
                              pinfo->num_domains);

pinfo->version = version;
handle->perf_ops = &perf_ops;
handle->perf_priv = pinfo;

return 0;
}

```

最终获取得到的OPP表如下

FTC664 (capacity: 576)				FTC310 (capacity: 1024)		
perf_value	power	trans_latency		perf_value	power	trans_latency
250MHz	79	1000		187.5MHz	1	1000
500MHz	197	1000		375MHz	9	1000
1000MHz	409	1000		750MHz	55	1000
2000MHz	839	1000		1500MHz	125	1000

## cpufreq初始化过程

```
// driver/base/cpu.c
struct bus_type cpu_subsys = {
    .name = "cpu",
    .dev_name = "cpu",
    .match = cpu_subsys_match,
    .online = cpu_subsys_online,
    .offline = cpu_subsys_offline,
};

// driver/cpufreq/cpufreq.c
// 指向当前使用的cpufreq_driver
static struct cpufreq_driver *cpufreq_driver;

// cpufreq subsys接口, 用来挂到CPU subsys总线上
static struct subsys_interface cpufreq_interface = {
    .name = "cpufreq",
    .subsys = &cpu_subsys,
    .add_dev = cpufreq_add_dev,
    .remove_dev = cpufreq_remove_dev,
};

// scmi cpufreq_driver结构体定义
static struct cpufreq_driver scmi_cpufreq_driver = {
    .name = "scmi",
    .flags = CPUFREQ_STICKY | CPUFREQ_HAVE_GOVERNOR_PER_POLICY |
             CPUFREQ_NEED_INITIAL_FREQ_CHECK,
    .verify = cpufreq_generic_frequency_table_verify,
    .attr = cpufreq_generic_attr,
    .target_index = scmi_cpufreq_set_target,
    .fast_switch = scmi_cpufreq_fast_switch,
    .get = scmi_cpufreq_get_rate,
    .init = scmi_cpufreq_init,
    .exit = scmi_cpufreq_exit,
    .ready = scmi_cpufreq_ready,
};

// 为每个cluster定义一个cpufreq_policy结构体, 对每个cluster上的CPU进行调频管理
// 其中又分别进行cpufreq_driver的初始化和governor的初始化
static DEFINE_PER_CPU(struct cpufreq_policy *, cpufreq_cpu_data);

// cpufreq驱动框架初始化过程, 整个过程都围绕着policy这个结构体进行, 逐步进行初始化
cpufreq_register_driver(&scmi_cpufreq_driver);
    subsys_interface_register(&cpufreq_interface);
        cpufreq_add_dev(dev, sif);
            cpufreq_online(cpu);
                // 初步初始化policy
                policy = cpufreq_policy_alloc(cpu);
                // 调用cpufreq_driver init接口, 完善policy结构体
                // 将opp表添加到对应的device, 通过dev_pm_opp_add接口
                // 生成频率表 freq_table
                cpufreq_driver->init(policy) -> scmi_cpufreq_init(policy)
                cpufreq_table_validate_and_sort(policy);
                // 创建/sys/device/system/cpu/cpux目录下的cpufreq符号链接
                add_cpu_dev_symlink();
                freq_qos_and_request();
                blocking_notifier_call_chain();
                // CPU进行频率调整, 使当前运行频率在频率表中
```

```
__cpufreq_driver_target();  
// 创建sys节点，/sys/device/system/cpu/cpufreq/policyx目录下的一些可选属性  
cpufreq_add_dev_interface(policy);  
cpufreq_stats_create_table(policy);  
list_add(&policy->polic_list, &cpufreq_policy_list);  
// 使用默认governor初始化policy  
cpufreq_init_policy();
```

## cpufreq\_governor的初始化过程

cpufreq governor的初始化过程，在cpufreq\_init\_policy(policy)中进行，这里以ondemand为例进行分析

```

// include/linux/cpufreq.h
struct cpufreq_governor {
    char    name[CPUFREQ_NAME_LEN];
    int     (*init)(struct cpufreq_policy *policy);
    void    (*exit)(struct cpufreq_policy *policy);
    int     (*start)(struct cpufreq_policy *policy);
    void    (*stop)(struct cpufreq_policy *policy);
    void    (*limits)(struct cpufreq_policy *policy);
    ssize_t (*show_setspeed)      (struct cpufreq_policy *policy,
                                   char *buf);
    int     (*store_setspeed)     (struct cpufreq_policy *policy,
                                   unsigned int freq);

    struct list_head    governor_list;
    struct module       *owner;
    u8                  flags;
};

/* Common Governor data across policies */
// 抽象出的governor调度器结构体
// drivers/cpufreq/cpufreq_governor.h
struct dbs_governor {
    struct cpufreq_governor gov;
    struct kobj_type kobj_type;
    /*
     * Common data for platforms that don't set
     * CPUFREQ_HAVE_GOVERNOR_PER_POLICY
     */
    struct dbs_data *gdbs_data;

    unsigned int (*gov_dbs_update)(struct cpufreq_policy *policy);
    struct policy_dbs_info *(*alloc)(void);
    void (*free)(struct policy_dbs_info *policy_dbs);
    int (*init)(struct dbs_data *dbs_data);
    void (*exit)(struct dbs_data *dbs_data);
    void (*start)(struct cpufreq_policy *policy);
};

// drivers/cpufreq/cpufreq_governor.h
// governor初始化宏
#define CPUFREQ_DBS_GOVERNOR_INITIALIZER(_name_) \
    { \
        .name = _name_, \
        .flags = CPUFREQ_GOV_DYNAMIC_SWITCHING, \
        .owner = THIS_MODULE, \
        .init = cpufreq_dbs_governor_init, \
        .exit = cpufreq_dbs_governor_exit, \
        .start = cpufreq_dbs_governor_start, \
        .stop = cpufreq_dbs_governor_stop, \
        .limits = cpufreq_dbs_governor_limits, \
    }

// ondemand governor定义
// driver/cpufreq/cpufreq_ondemand.c
static struct dbs_governor od_dbs_gov = {
    .gov = CPUFREQ_DBS_GOVERNOR_INITIALIZER("ondemand"),
    .kobj_type = { .default_attrs = od_attributes },
    .gov_dbs_update = od_dbs_update,
    .alloc = od_alloc,
    .free = od_free,
};

```

```

        .init = od_init,
        .exit = od_exit,
        .start = od_start,
};

// 初始化governor
// 该函数会在governor模块驱动的入口函数调用
// 只要编译该模块，就会注册到cpufreq framework中
#define CPU_FREQ_GOV_ONDEMAND (od_dbs_gov.gov)
cpufreq_governor_init(CPU_FREQ_GOV_ONDEMAND);

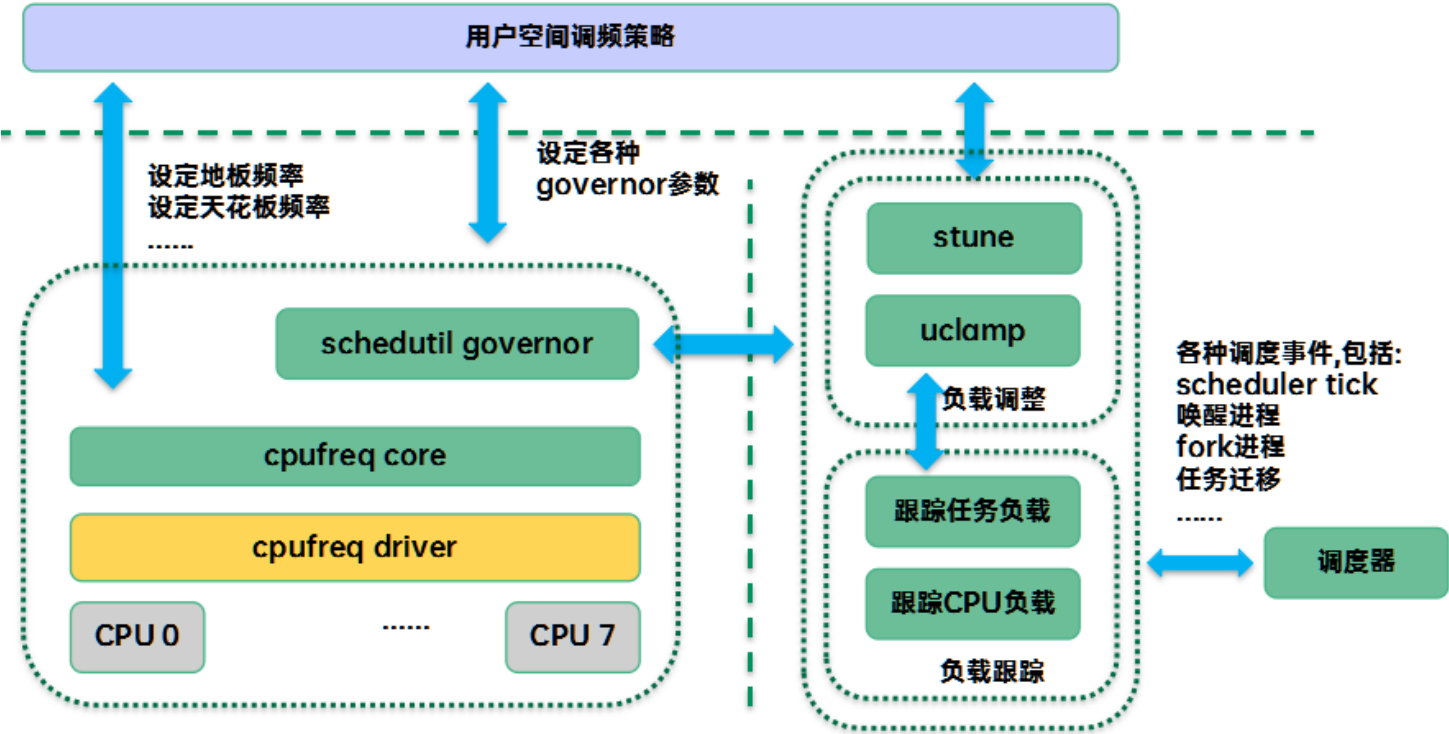
#define cpufreq_governor_init(__governor) \
static int __init __governor##_init(void) \
{ \
    return cpufreq_register_governor(&__governor); \
} \
core_initcall(__governor##_init)

// 在cpufreq_online()中调用默认governor对policy进行完善，启动当前governor
cpufreq_init_policy(policy);
gov = get_governor(default_governor);
// 设置新的governor
cpufreq_set_policy(policy, gov, pol);
cpufreq_init_governor(policy);
    policy->governor->init(); -> cpufreq_dbs_governor_init(policy);
    gov->init(dbs_data); -> odinit(dbs_data);
cpufreq_start_governor(policy);
    policy->governor->start(); -> cpufreq_dbs_governor_start(policy);
    // 设置governor回调函数，以ondemand为例
    gov_set_update_util(policy_dbs, sampling_rate);
    cpufreq_add_update_util_hook(cpu, &cddb->update_util,
                                dbs_update_util_handler);

```

启动governor中比较重要的是设置调频回调函数,该函数是真正调频时计算合适频率的函数

# schedutil调节器



sugov作为一种内核调频策略模块，它主要是根据当前CPU的利用率进行调频。因此，sugov会注册一个callback函数（sugov\_update\_shared/sugov\_update\_single）到调度器负载跟踪模块，当CPU util发生变化的时候就会调用该callback函数，检查一下当前CPU频率是否和当前的CPU util匹配，如果不匹配，那么就进行升频或者降频。

sugov\_tunables 结构体，用来描述sugov的可调参数

```
// sugov_tunables结构体
struct sugov_tunables {
    struct gov_attr_set attr_set;
    // 目前sugov只有这一个可调参数，该参数用来限制连续调频的间隔时间
    unsigned int rate_limit_us;
};
```

sugov\_policy 结构体，sugov为每个cluster构建了该数据结构，记录每个cluster的调频数据信息

// sugov\_policy结构体，为每个簇构建了该数据结构，记录每个簇的调频数据信息

```
struct sugov_policy {
    // 指向cpufreq framework层的policy对象
    struct cpufreq_policy *policy;
    // sugov的可调参数
    struct sugov_tunables *tunables;
    struct list_head tunables_hook;

    raw_spinlock_t update_lock; /* For shared policies */
    // 记录上次进行频率调整的时间点
    u64 last_freq_update_time;
    // 最小调频时间间隔
    s64 freq_update_delay_ns;
    // 下一个需要调整到的频率值，回调函数主要是计算这个参数
    unsigned int next_freq;
    // 根据CPU util计算出来的原始频率，在频率表中向上找最接近的频率进行调整
    unsigned int cached_raw_freq;

    /* The next fields are only needed if fast switch cannot be used: */
    struct irq_work irq_work;
    struct kthread_work work;
    struct mutex work_lock;
    struct kthread_worker worker;
    struct task_struct *thread;
    bool work_in_progress;

    bool limits_changed;
    bool need_freq_update;
};
```

sugov\_cpu 结构体，sugov为每个cpu构建了该数据结构，记录per-cpu的调频数据信息

```
struct sugov_cpu {
    // 保存了cpu util变化后的回调函数
    struct update_util_data update_util;
    // 该sugov_cpu对应的sugov_policy对象
    struct sugov_policy *sg_policy;
    // 对应的CPU id
    unsigned int cpu;

    bool iowait_boost_pending;
    unsigned int iowait_boost;
    // 上一次cpu负载变化驱动调频的时间点，util更新的时间点
    u64 last_update;

    unsigned long bw_dl;
    // 该CPU的最大算力，即最大utilities，归一化到1024
    unsigned long max;

    /* The field below is for single-CPU policies only: */
#ifdef CONFIG_NO_HZ_COMMON
    unsigned long saved_idle_calls;
#endif
};
```



sugov初始化过程和ondemand初始化过程相似，当内核设定默认governor为sugov时，在 `cpufreq_init_governor(policy);` 中会调用 `sugov_init()` 初始化sugov，然后调用 `sugov_start()` 设置调频回调函数，每当CPU利用率发生变化的时候，调度器都会调用 `cpufreq_update_util()` 通知sugov，在 `cpufreq_update_util()` 被调用时，即任务调度后CPU当前的util发生变化，会调用sugov的回调函数进行调频，`sugov_update_shared()` 当一个簇中有多个CPU调用该回调，遍历簇上的CPU找到当前最大util的CPU，然后根据该util映射到频率；`sugov_update_single()` 即一个簇上单个CPU的情况直接根据该CPU的 `util` 计算频率

调度事件的发生还是非常密集的，特别是在重载的情况下，很多任务可能执行若干个us就切换出去了。如果每次都计算CPU util看看是否需要调整频率，那么本身sugov就给系统带来较重的负荷，因此并非每次调频时机都会真正执行调频检查，sugov设置了一个最小调频间隔，小于这个间隔的调频请求会被过滤掉。

## schedutil频率计算过程

```
// sugov_start会遍历该sugov policy (cluster) 中的所有cpu
// 调用cpufreq_add_update_util_hook为sugov cpu注册调频回调函数，代码逻辑如下：
static int sugov_start(struct cpufreq_policy *policy)
{
    ...
    for_each_cpu(cpu, policy->cpus) {
        struct sugov_cpu *sg_cpu = &per_cpu(sugov_cpu, cpu);
        // 设置governor 计算回调函数，cpufreq_update_util()被调用时
        // 即任务调度后CPU当前的util发生变化，会调用sugov的回调函数进行调频计算
        cpufreq_add_update_util_hook(cpu, &sg_cpu->update_util,
                                     policy_is_shared(policy) ?
                                     sugov_update_shared :
                                     sugov_update_single);
    }
    ...
}

// schedutil频率计算过程
sugov_update_single();
// 调频最小间隔时间检查，小于设定时间，直接返回
sugov_should_update_freq();
util = sugov_get_util(sg_cpu);
    schedutil_cpu_util(sg_cpu->cpu, util, max, FREQUENCY_UTIL, NULL);
// 根据当前CPU的util映射到具体的频率上
next_f = get_next_freq(sg_policy, util, max);
// 调用cpufreq_driver进行调频
sugov_deferred_update(sg_policy, time, next_f);
    __cpufreq_driver_target()

// 计算cpu当前的utility
unsigned long schedutil_cpu_util(int cpu, unsigned long util_cfs,
                                unsigned long max, enum schedutil_type type,
                                struct task_struct *p)
{
    unsigned long dl_util, util, irq;
    struct rq *rq = cpu_rq(cpu);

    if (!uclamp_is_used() &&
        type == FREQUENCY_UTIL && rt_rq_is_runnable(&rq->rt)) {
        return max;
    }

    // 如果CPU处理了过多的中断服务函数，irq负载已经高过CPU最大算力，直接返回最大算力
    irq = cpu_util_irq(rq);
    if (unlikely(irq >= max))
        return max;

    // 累加了cfs和rt任务的utility
    util = util_cfs + cpu_util_rt(rq);
    if (type == FREQUENCY_UTIL)
        util = uclamp_rq_util_with(rq, util, p);

    dl_util = cpu_util_dl(rq);

    if (util + dl_util >= max)
        return max;
```

```

/*
 * OTOH, for energy computation we need the estimated running time, so
 * include util_dl and ignore dl_bw.
 */
if (type == ENERGY_UTIL)
    util += dl_util;

/*
 * There is still idle time; further improve the number by using the
 * irq metric. Because IRQ/steal time is hidden from the task clock we
 * need to scale the task numbers:
 *
 *
 *          max - irq
 * U' = irq + ----- * U
 *          max
 */
// irq会偷走一部分的cpu算力, 从而让其capacity没有那么大。
// 这里通过scale_irq_capacity对任务的utility进行调整
util = scale_irq_capacity(util, irq, max);
util += irq;

/*
 * Bandwidth required by DEADLINE must always be granted while, for
 * FAIR and RT, we use blocked utilization of IDLE CPUs as a mechanism
 * to gracefully reduce the frequency when no tasks show up for longer
 * periods of time.
 *
 * Ideally we would like to set bw_dl as min/guaranteed freq and util +
 * bw_dl as requested freq. However, cpufreq is not yet ready for such
 * an interface. So, we only do the latter for now.
 */
if (type == FREQUENCY_UTIL)
    util += cpu_bw_dl(rq);

return min(max, util);
}

// 根据当前CPU计算的util映射对应频率
static unsigned int get_next_freq(struct sugov_policy *sg_policy,
                                unsigned long util, unsigned long max)
{
    struct cpufreq_policy *policy = sg_policy->policy;
    // 先取得当前CPU的最大频率
    unsigned int freq = arch_scale_freq_invariant() ?
                        policy->cpuinfo.max_freq : policy->cur;
    // 计算当前util对应频率, 计算公式: freq = (1.25) * freq * util / max
    // 这里冗余了25%的算力余量
    freq = map_util_freq(util, freq, max);

    // 若计算出的freq和上次缓存的一样, 则实际调整的next_freq计算后肯定也是一样的, 直接返回
    // 上次记录的频率值
    if (freq == sg_policy->cached_raw_freq && !sg_policy->need_freq_update)
        return sg_policy->next_freq;

    sg_policy->cached_raw_freq = freq;

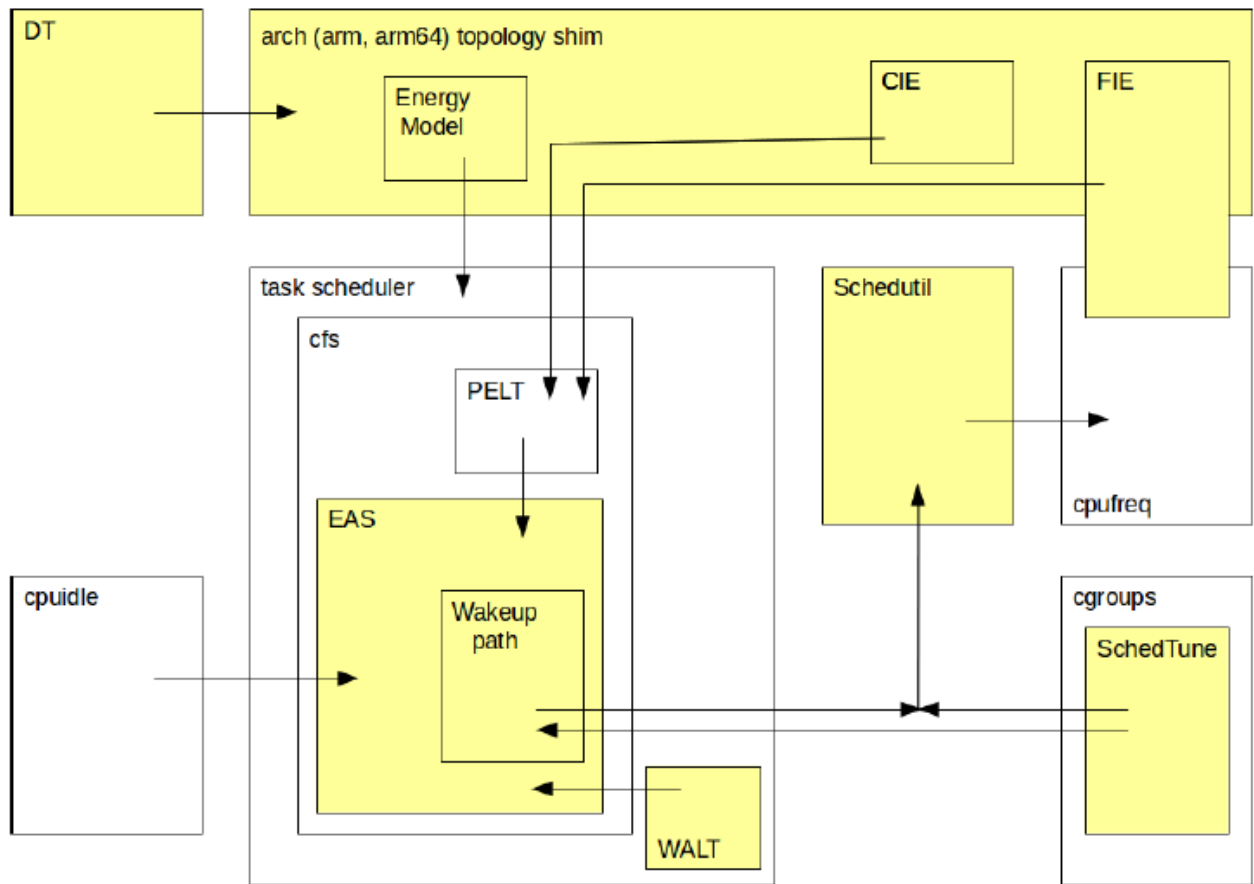
    // 根据当前算的freq, 在CPU频率表上查找对应的频率
    freq = cpufreq_driver_resolve_freq(policy, freq);

```

```
    return freq;
}
```

## EAS能源感知调度

EAS整体框架



**Figure 1 EAS building blocks in relation to Linux task scheduler, cgroups subsystem and related power management subsystems**

完全公平调度（Completely Fair Scheduler CFS）实现了面向吞吐量的任务调度策略，EAS为这个调度器添加了一个基于能耗的调度策略，在优化CPU算力冗余的同时实现了节能，EAS在系统中、低度负载情况下工作，CFS在系统满负载情况下工作。

EAS在CPU调度领域，在为任务选核起作用，目的是保证性能的情况下尽可能节省功耗，EAS涉及内核的几个子系统（任务调度、能源管理、CPU动态调频），EAS代码主要位于 `kernel/sched/fair.c`，能源感知的任务调度需要调度器评估各个任务在CPU上运行带来的能耗影响

EAS负载跟踪有两种模式，一种是“每实体负载跟踪（Per\_Entity Load Track）”，通常用于负载跟踪，然后该信息用于确定频率以及如何在CPU上委派任务，另一种是“窗口辅助的负载跟踪（Window-Assisted Load Tracking）”，WALT更具有突发性，而PELT试图让频率保持连贯性，负载跟踪器实际上并不影响CPU频率，它只是告诉系统CPU使用率是多少

EAS全局控制开关 `/proc/sys/kernel/sched_energy_aware`

CPU算力归一化过程

当前，Linux无法凭自身算出CPU算力，因此必须要有把这个信息传递给Linux的方式，它是从 capacity-dmips-mhz CPU 设备树binding中衍生计算出来的

归一化CPU capacity， topology\_normalize\_cpu\_scale() 定义在 drivers/base/arch\_topology() ，这个capacity在 schedutil调度中被 sugov\_get\_util() 函数读取

topology\_normalize\_cpu\_scale() 在CPU初始化 parse\_dt\_topology() 中被调用，capacity归一化的前提条件是需要 在设备树中CPU节点设置 capacity-dmips-mhz 属性，该属性表示不同CPU的计算能力，内核读取该属性设置CPU 的 raw\_capacity 为 capacity-dmips-mhz ，参考内核文 档 Documentation/devicetree/bindings/arm/cpu-capacity.txt

ARM推荐的测试CPU的性能工具：Dhrystone 2.1以上版本，可以通过单核跑分成绩作为 capacity-dmips-mhz 属性的参考，DMIPS： Dhrystone Million Instructions executed Per Second，表示了Dhrystone这样一种测试方法下的MIPS，Dhrystone是一种整数运算测试程序。MIPS/MHz，就是说每MHz频率能产生多大的MIPS，CPU性能通常由每秒百万指令（Millions of Instructions Per Second，MIPS）表示，设备树里表示为dmips/mhz

CPU算力归一化公式，并不是简单的将capacity-dmips-mhz归一化到capacity，CPU的频率也参与到了计算中

$$capacity = (own(capacity-dmips-mhz) * own(max\_freq)) / (max(capacity-dmips-mhz) * max(max\_freq)) * 1024$$

根据测试部测试的E2000QCPU单核性能数据，E2000Q的 capacity-dmips-mhz 属性值可以设置为如下，放大1000倍：

dhrystone	Dhrystones/s	N/A	19920318	7530120
	DMIPS/MHz	N/A	5.66	2.85

```
// 小核
cpu_l0: cpu@0 {
    ...
    capacity-dmips-mhz = <2850>;
    ...
};
// 大核
cpu_b0: cpu@0 {
    ...
    capacity-dmips-mhz = <5660>;
    ...
};
```

实际经过CPU算力归一化到1024之后，对应的小核CPU算力为386，大核为1024

EAS代码相关结构体

perf\_domain结构表示一个CPU性能域，perf\_domain和cpufreq\_policy是一一对应的，性能域之间形成链，链表头存放在 root\_domian中

```

// kernel/sched/sched.h
// perf_comain 结构表示一个CPU性能域，perf_domain和cpufreq_policy是一一对应的
struct perf_domain {
    struct em_perf_domain *em_pd;
    struct perf_domain *next;
    struct rcu_head rcu;
};

struct root_domain {
    atomic_t          refcount;
    atomic_t          rto_count;
    struct rcu_head    rcu;
    cpumask_var_t      span;
    cpumask_var_t      online;

    // 该root_domain是否处于overload状态
    int                overload;

    // 该root_domain是否处于overutilized状态
    int                overutilized;

    /*
     * The bit corresponding to a CPU gets set here if such CPU has more
     * than one runnable -deadline task (as it is below for RT tasks).
     */
    cpumask_var_t      dlo_mask;
    atomic_t           dlo_count;
    struct dl_bw        dl_bw;
    struct cpudl        cpudl;

#ifdef HAVE_RT_PUSH_IPI
    /*
     * For IPI pull requests, loop across the rto_mask.
     */
    struct irq_work      rto_push_work;
    raw_spinlock_t       rto_lock;
    /* These are only updated and read within rto_lock */
    int                  rto_loop;
    int                  rto_cpu;
    /* These atomics are updated outside of a lock */
    atomic_t             rto_loop_next;
    atomic_t             rto_loop_start;
#endif

    /*
     * The "RT overload" flag: it gets set if a CPU has more than
     * one runnable RT task.
     */
    cpumask_var_t        rto_mask;
    struct cpupri        cpupri;

    // 系统中算力最大的CPU的算力
    unsigned long        max_cpu_capacity;

    /*
     * NULL-terminated list of performance domains intersecting with the
     * CPUs of the rd. Protected by RCU.
     */
    // perf_domain单链表的表头
    struct perf_domain __rcu *pd;

```

```

};

// include/linux/energy_model.h
struct em_perf_state {
    unsigned long frequency;           // CPU频点, 单位KHz
    unsigned long power;               // 此频点下的功耗
    unsigned long cost;                // 此频点下的成本系数, 等于 power * max_freq / freq
};

struct em_perf_domain {
    struct em_perf_state *table;       // CPU频点表
    int nr_perf_states;                // 频点表中元素的个数
    unsigned long cpus[];              // 此性能域中包括哪些CPU
};

```

E2000Q 5.10内核, perf\_domain\_debug 打印信息

```

[ 2.574534] root_domain 0-3: pd3:{ cpus=3 nr_pstate=4 }
[ 2.574540] freq: 250000, power: 79, cost: 632
[ 2.579072] freq: 500000, power: 197, cost: 788
[ 2.583690] freq: 1000000, power: 409, cost: 818
[ 2.588390] freq: 2000000, power: 839, cost: 839
[ 2.593094] pd2:{ cpus=2 nr_pstate=4 }
[ 2.593096] freq: 250000, power: 79, cost: 632
[ 2.601445] freq: 500000, power: 197, cost: 788
[ 2.606054] freq: 1000000, power: 409, cost: 818
[ 2.610749] freq: 2000000, power: 839, cost: 839
[ 2.615445] pd0:{ cpus=0-1 nr_pstate=4 }
[ 2.615447] freq: 187500, power: 1, cost: 8
[ 2.623709] freq: 375000, power: 9, cost: 36
[ 2.628058] freq: 750000, power: 55, cost: 110
[ 2.632579] freq: 1500000, power: 125, cost: 125

```

root\_domain的overload和overutilized说明:

- 对于一个 CPU 而言, 其处于 overload 状态则说明其 rq 上有大于等于2个任务, 或者虽然只有一个任务, 但是是 misfit task
- 对于一个 CPU 而言, 其处于 overutilized 状态说明该 cpu 的 utility 超过其 capacity (缺省预留20%的算力, 另外, 这里的 capacity 是用于cfs任务的算力)
- 对于 root domain, overload 表示至少有一个 cpu 处于 overload 状态。overutilized 表示至少有一个 cpu 处于 overutilized 状态
- overutilized 状态非常重要, 它决定了调度器是否启用EAS, 只有在系统没有 overutilized 的情况下EAS才会生效。overload和newidle balance的频次控制相关, 当系统在overload的情况下, newidle balance才会启动进行均衡。

## EAS能量计算方法

CPU在某个performance state(ps)下的计算能力:

$ps \rightarrow cap = ps \rightarrow freq * scale\_cpu / cpu\_max\_freq$  (1)

CPU在该频点performace state(ps)下的能量消耗:

$cpu\_nrg = ps \rightarrow power * cpu\_util / ps \rightarrow cap$  (2)

结合(1) (2)可以得出CPU在该ps下的能量消耗

$$\text{cpu\_nrg} = \text{ps} \rightarrow \text{power} * \text{cpu\_max\_freq} * \text{cpu\_util} / \text{ps} \rightarrow \text{freq} * \text{scale\_cpu} \quad (3)$$

其中  $\text{ps} \rightarrow \text{power} * \text{cpu\_max\_freq} / \text{ps} \rightarrow \text{freq}$  是一个固定数据存放在频点表的cost成员中

一个pd内的CPU，拥有相同的cost，所以一个pd内所有CPU的能量消耗可以表示为

$$\text{pd\_nrg} = \text{ps} \rightarrow \text{cost} * \text{sum}(\text{cpu\_util}) / \text{scale\_cpu}$$

### **EAS的调度过程**

在任务被重新唤醒或者fork新建时，会通过 `select_task_rq_fair()` 将任务进行balance，达到充分利用CPU的目的。

在 `select_task_rq_fair()`，若任务是被重新唤醒就会调用 `find_energy_efficient_cpu()` 进行选核执行



```

/*
 * compute_energy(): Estimates the energy that @pd would consume if @p was
 * migrated to @dst_cpu. compute_energy() predicts what will be the utilization
 * landscape of @pd's CPUs after the task migration, and uses the Energy Model
 * to compute what would be the energy if we decided to actually migrate that
 * task.
 */
// 计算任务迁移到dst_cpu后，整个pd，即此cluster的能量
static long
compute_energy(struct task_struct *p, int dst_cpu, struct perf_domain *pd)
{
    struct cpumask *pd_mask = perf_domain_span(pd);
    // 获取该CPU的算力
    unsigned long cpu_cap = arch_scale_cpu_capacity(cpumask_first(pd_mask));
    unsigned long max_util = 0, sum_util = 0;
    int cpu;

    // 对此pd中每个online cpu都执行计算
    for_each_cpu_and(cpu, pd_mask, cpu_online_mask) {
        unsigned long cpu_util, util_cfs = cpu_util_next(cpu, p, dst_cpu);
        struct task_struct *tsk = cpu == dst_cpu ? p : NULL;

        // 返回该CPU下cfs+irq+rt+dl使用掉的CPU算力总和
        sum_util += schedutil_cpu_util(cpu, util_cfs, cpu_cap,
                                       ENERGY_UTIL, NULL);

        cpu_util = schedutil_cpu_util(cpu, util_cfs, cpu_cap,
                                       FREQUENCY_UTIL, tsk);
        max_util = max(max_util, cpu_util);
    }

    // 计算该pd下所有CPU的功耗和
    return em_cpu_energy(pd->em_pd, max_util, sum_util);
}

// 计算该pd下所有CPU的功耗和
static inline unsigned long em_cpu_energy(struct em_perf_domain *pd,
                                          unsigned long max_util, unsigned long sum_util)
{
    unsigned long freq, scale_cpu;
    struct em_perf_state *ps;
    int i, cpu;

    /*
     * In order to predict the performance state, map the utilization of
     * the most utilized CPU of the performance domain to a requested
     * frequency, like schedutil.
     */
    cpu = cpumask_first(to_cpumask(pd->cpus));
    scale_cpu = arch_scale_cpu_capacity(cpu);
    ps = &pd->table[pd->nr_perf_states - 1];
    freq = map_util_freq(max_util, ps->frequency, scale_cpu);

    /*
     * Find the lowest performance state of the Energy Model above the
     * requested frequency.
     */
    for (i = 0; i < pd->nr_perf_states; i++) {
        ps = &pd->table[i];

```

```

        if (ps->frequency >= freq)
            break;
    }

/*
 * The capacity of a CPU in the domain at the performance state (ps)
 * can be computed as:
 *
 *      ps->freq * scale_cpu
 *      ps->cap = -----          (1)
 *      cpu_max_freq
 *
 * So, ignoring the costs of idle states (which are not available in
 * the EM), the energy consumed by this CPU at that performance state
 * is estimated as:
 *
 *      ps->power * cpu_util
 *      cpu_nrg = -----          (2)
 *      ps->cap
 *
 * since 'cpu_util / ps->cap' represents its percentage of busy time.
 *
 * NOTE: Although the result of this computation actually is in
 * units of power, it can be manipulated as an energy value
 * over a scheduling period, since it is assumed to be
 * constant during that interval.
 *
 * By injecting (1) in (2), 'cpu_nrg' can be re-expressed as a product
 * of two terms:
 *
 *      ps->power * cpu_max_freq   cpu_util
 *      cpu_nrg = ----- * -----          (3)
 *      ps->freq                   scale_cpu
 *
 * The first term is static, and is stored in the em_perf_state struct
 * as 'ps->cost'.
 *
 * Since all CPUs of the domain have the same micro-architecture, they
 * share the same 'ps->cost', and the same CPU capacity. Hence, the
 * total energy of the domain (which is the simple sum of the energy of
 * all of its CPUs) can be factorized as:
 *
 *      ps->cost * \Sum cpu_util
 *      pd_nrg = -----          (4)
 *      scale_cpu
 */
return ps->cost * sum_util / scale_cpu;
}

```

// 寻找工作能耗最低的CPU

```

static int find_energy_efficient_cpu(struct task_struct *p, int prev_cpu)
{
    unsigned long prev_delta = ULONG_MAX, best_delta = ULONG_MAX;
    struct root_domain *rd = cpu_rq(smp_processor_id())->rd;
    unsigned long cpu_cap, util, base_energy = 0;
    int cpu, best_energy_cpu = prev_cpu;
    struct sched_domain *sd;
    struct perf_domain *pd;

```

```

rcu_read_lock();
// 从rd取pd的指针
pd = rcu_dereference(rd->pd);
if (!pd || READ_ONCE(rd->overutilized))
    goto fail;

/*
 * Energy-aware wake-up happens on the lowest sched_domain starting
 * from sd_asym_cpucapacity spanning over this_cpu and prev_cpu.
 */
sd = rcu_dereference(*this_cpu_ptr(&sd_asym_cpucapacity));
while (sd && !cpumask_test_cpu(prev_cpu, sched_domain_span(sd)))
    sd = sd->parent;
if (!sd)
    goto fail;

sync_entity_load_avg(&p->se);
if (!task_util_est(p))
    goto unlock;

// 遍历整个pd链表，计算p在不同pd下的能耗
for (; pd; pd = pd->next) {
    unsigned long cur_delta, spare_cap, max_spare_cap = 0;
    unsigned long base_energy_pd;
    int max_spare_cap_cpu = -1;

    /* Compute the 'base' energy of the pd, without @p */
    // 计算不包括p的情况下此pd当前的energy
    base_energy_pd = compute_energy(p, -1, pd);
    // 不包括p的情况下系统的总energy
    base_energy += base_energy_pd;

    // 遍历整个pd中的CPU，计算p放在该CPU上的功耗
    for_each_cpu_and(cpu, perf_domain_span(pd), sched_domain_span(sd)) {
        if (!cpumask_test_cpu(cpu, p->cpus_ptr))
            continue;
        // 计算p放到此CPU后该CPU总共消耗的算力
        util = cpu_util_next(cpu, p, cpu);
        cpu_cap = capacity_of(cpu);
        spare_cap = cpu_cap;
        // 计算p放到此CPU后剩余的算力
        lsub_positive(&spare_cap, util);

        /*
         * Skip CPUs that cannot satisfy the capacity request.
         * IOW, placing the task there would make the CPU
         * overutilized. Take uclamp into account to see how
         * much capacity we can get out of the CPU; this is
         * aligned with schedutil_cpu_util().
         */
        util = uclamp_rq_util_with(cpu_rq(cpu), util, p);
        // CPU需要保留20%左右的算力，不满足需求后进行下一个CPU的探测
        if (!fits_capacity(util, cpu_cap))
            continue;

        /* Always use prev_cpu as a candidate. */
        // 若对比的这个CPU就是任务之前运行的CPU
        if (cpu == prev_cpu) {
            // 计算p放在该cpu后整个pd的能量消耗

```

```

        prev_delta = compute_energy(p, prev_cpu, pd);
        // 计算p放在该CPU后整个pd增加的能量消耗
        prev_delta -= base_energy_pd;
        // 更新best_delta, 取最优能耗
        best_delta = min(best_delta, prev_delta);
    }

    /*
     * Find the CPU with the maximum spare capacity in
     * the performance domain
     */
    // 记录p放上去后剩余算力最大的CPU和最大的剩余算力
    if (spare_cap > max_spare_cap) {
        max_spare_cap = spare_cap;
        max_spare_cap_cpu = cpu;
    }
}

/* Evaluate the energy impact of using this CPU. */
// 同一个簇上的CPU取最大余量算力的那个CPU与其他簇的CPU做能量消耗对比
if (max_spare_cap_cpu >= 0 && max_spare_cap_cpu != prev_cpu) {
    // 计算p放在算力剩余最大的CPU后整个pd的能量消耗
    cur_delta = compute_energy(p, max_spare_cap_cpu, pd);
    // 计算能量消耗增量
    cur_delta -= base_energy_pd;
    // 如果当前能量增量优于p放在prev_cpu运行的能量消耗, 则取该cpu运行p
    if (cur_delta < best_delta) {
        best_delta = cur_delta;
        best_energy_cpu = max_spare_cap_cpu;
    }
}

}

unlock:
    rcu_read_unlock();

    /*
     * Pick the best CPU if prev_cpu cannot be used, or if it saves at
     * least 6% of the energy used by prev_cpu.
     */
    // 若prev_cpu找不到, 就直接返回最优能耗cpu
    if (prev_delta == ULONG_MAX)
        return best_energy_cpu;

    // 若最优能耗比放在prev_cpu上运行的能耗还要低6.25%以上, 则取最优能耗cpu
    if ((prev_delta - best_delta) > ((prev_delta + base_energy) >> 4))
        return best_energy_cpu;

    // 否则不做改变, 直接使用prev_cpu运行p
    return prev_cpu;

fail:
    rcu_read_unlock();

    return -1;
}

```